

# Precise reference counting for lazy functional languages with interprocedural analysis

Martin Fredin  
`fredinm@chalmers.se`

2023

## Abstract

Precise reference counting is a technique by Reinking et al. that uses ownership to deallocate objects as soon as possible. The algorithm is called Perceus, and as of this writing, it has only been implemented for eager functional languages. This paper describes the implementation of a new lazy compiler back-end for the Agda programming language with precise reference counting. The compiler uses Boquist and Johnsson’s intermediate language GRIN to compile lazy programs. GRIN uses interprocedural analysis to inline the evaluation of suspended computations. We extend GRIN with a variant of Perceus, and demonstrate the applicability of combining lazy functional programming with precise reference counting by developing a GRIN interpreter and an LLVM code generator.

## 1 Introduction

Reference counting (Collins, 1960) is a memory management technique which can detect and free resources as soon as they are no longer needed, allowing memory reuse and destructive updates. Common reference counting algorithms are easy to implement; each allocation contains an extra field which tracks the number of references to an object. When the reference count reaches zero, the heap space occupied by the object is reclaimed. The number of references to a object is updated by interleaved reference counting operations (`dup` and `drop`), which increments and decrements the reference count at runtime. As a result of the interleaved collection strategy, memory usage remain low and the throughput is continuous throughout

the computation<sup>1</sup> (Jones & Lins, 1996). Still, tracing garbage collectors are usually favored over reference counting implementations due to cheaper allocations, higher throughput, and the ability to collect cyclic data structures.

Reinking et al. (2021) reexamine reference counting with a new approach, utilizing static guarantees to make the algorithm precise so objects can be deallocated as soon as possible. They present a formalized algorithm called Perceus, which ensures preciseness. Perceus is implemented in the functional language Koka, along with optimizations for reducing reference counting operations and reusing memory. This, builds upon previous work by Ullrich and de Moura (2021) in the Lean programming language and theorem prover.

Both Koka and Lean are, however, eagerly evaluated. Lazy languages pose an extra challenge for compiler writers because of their unintuitive control flow. In this paper, we adapt the Perceus algorithm to a new lazy compiler back-end for the Agda programming language and proof assistant. As a first step, we transform Agda into an intermediate language called GRIN (Johnsson, 1991).

## 2 Graph Reduction Intermediate Notation

In 1991, Johnsson presented the Graph Reduction Intermediate Notation (GRIN) as an imperative version of the G-machine (Johnsson, 1984), where

---

<sup>1</sup>Reference counted programs may introduce pauses similar to tracing garbage collectors. For example, when decrementing a long linked list all at once.

lexically scoped variables are stored in registers instead of the stack. Later, GRIN was reformulated with a more "functional flavor" (Boquist, 1995). In this project, we introduce an additional variant of GRIN adapted for the internal representation of Agda and precise reference counting.

The syntax of our variant are shown in Figure 1. There are 4 kinds of constructs; terms, values, lambda patterns, and case patterns. All syntactically correct expressions are not valid. For example, the value at the function position of an application must be either a top-level function (*def*) or a primitive (*prim*). It cannot be a variable because there are no indirect function calls in GRIN. Likewise, top-level functions cannot be passed as arguments because GRIN is a first-order language.

## 2.1 Code generation

The current implementation of GRIN only compiles a subset of Agda which is lambda lifted, first order, and monomorphic.

```

downFrom : ℕ → List ℕ
downFrom zero = []
downFrom (suc n) = n :: downFrom n

sum : List ℕ → ℕ
sum [] = 0
sum (x :: xs) = x + sum xs

main = sum (downFrom 100)

```

Our back-end starts by converting the program into the treeless syntax. The treeless syntax has explicit case expressions which use A-normal-form, meaning that the scrutinee is always a variable, and the alternatives cannot be nested or overlap. Following is the treeless representation of `downFrom`.

```

downFrom x7 = case x7 of
  0 → []
  _ → let x5 = 1
        x4 = _-_ x7 x5
        x3 = downFrom x4 in
        _::_ x4 x3

```

The implementation uses de Bruijn indices to represent variables, but this paper uses variable names to make it more readable. During this phase, we also

transform the program so applications only take variables as operands.

GRIN is very similar to the treeless syntax, but instead of let-expressions we use the builtin state monad to bind variables and sequence operations. The monadic operations are **unit**, **store**, **fetch**, and **update**. The bind operator ";" is infix and right-associative. We can translate "let x = val in foo x" lazily by allocating the value and passing the pointer as an argument to the function: "store val ; λ x → foo x". Here, *val* must be a constant node value. A constant node is a tag followed by a sequence of arguments. We can pattern match on a tag with the case expression to determine the kind. Tags are prefixed with either a "C" if it is a constructor value or an "F" for suspended function applications. The node arguments are usually pointers to other heap-allocated nodes, but they can also be unboxed values. For example, the boxed integer tag Cnat accepts one unboxed integer.

```

downFrom x7 =
  eval x7 ; λ Cnat x6 →
  case x6 of
    0 → unit (C[])
    _ →
      store (Cnat 1) ; λ x5 →
      store (F_ _ x7 x5) ; λ x4 →
      store (FdownFrom x4) ; λ x3 →
      unit (C_::_ x4 x3)

```

## 2.2 Analysis and transformations

The most important GRIN transformation is *eval inlining*. `eval` is a normal GRIN function which forces suspended computations to its weak head normal form. There are no suspended computations in GRIN and all values are weak head normal form. Thus, "suspended computation" and "weak head normal form" should always refer to the corresponding representation in the source program. In the above function, "eval x7 ; λ Cnat x6 → ..." evaluates the value at the pointer (x7) to a boxed integer Cnat x6.

Eval inlining generates a specialized `eval` function for each call site. To evaluate a suspended computation, we load the node from the heap us-

$term$	$::=$	$term ; \lambda lpat \rightarrow term$	binding
		<b>case</b> $val$ <b>of</b> $term \{calt\}^*$	case
		$val \{val\}^*$	application
		<b>unit</b> $val$	return value
		<b>store</b> $val$	allocate new heap node
		<b>fetch</b> $\{tag\} n \{i\}$	load heap node
		<b>update</b> $\{tag\} n \{i\} val$	overwrite heap node
		<b>unreachable</b>	unreachable
$val$	$::=$	$tag \{val\}^*$	constant node
		$n \{val\}^*$	variable node
		$tag$	single tag
		$()$	empty
		$lit$	literal
		$n$	variable (de Bruijn index)
		$def$	function definition
		$prim$	primitive definition
$lpat$	$::=$	$tag \{x\}^*$	constant node pattern
		$x \{x\}^*$	variable node pattern
		$()$	empty pattern
		$x$	variable pattern
$cpat$	$::=$	$tag \{x\}^*$	constant node pattern
		$tag$	single tag pattern
		$lit$	literal pattern
$\{\dots\}$			means 0 or 1 times
$\{\dots\}^*$			means 0 or more times

**Figure 1:** *GRIN syntax.*

ing **fetch**. Then, we pattern match on the possible nodes. Constructor nodes are already in weak head normal form so they are left unchanged. Function nodes need to be evaluated by applying the arguments to the corresponding function. Finally, the heap is updated with the evaluated value.

```

downFrom x7 =
  (fetch x7 ; λ x34 →
    (case x34 of
      Cnat x36 → unit (Cnat x36)
      F_ _ x37 x38 → _ _ x37 x38
    ) ; λ x35 →
    update x7 x34 ; λ () →
    unit x35
  ) ; λ Cnat x6 →
  case x6 of
    0 → unit (C[])
    _ →
      store (Cnat 1) ; λ x5 →
      store (F_ _ x7 x5) ; λ x4 →
      store (FdownFrom x4) ; λ x3 →
      unit (C_::_ x4 x3)

```

Eval inlining require a set of possible nodes for each abstract heap location. The set needs to be relatively small, or otherwise an excessive amount of code will be generated. We use the *heap points-to* analysis (Johnsson, 1991). The analysis is interprocedural, meaning that multiple functions need to be analyzed together. We will not go into detail about the algorithm, as it is thoroughly described in (Boquist & Johnsson, 1996). Instead, this paper will only provide a general intuition of the algorithm. Consider the inlined evaluation in `downFrom`. There are two tags in the case expression: `F_ _` and `Cnat`. `F_ _` comes from the suspended recursive call inside `downFrom`, and `Cnat` is from the **update** operation and the suspended call to `downFrom` in the main function.

```

main =
  store (Cnat 100) ; λ x20 →
  store (FdownFrom x20) ; λ x19 →
  sum x19 ; λ Cnat x18 →
  printf x18

```

Boquist’s thesis presents 24 transformations divided into two groups: simplifying transformations and optimizing transformations (Boquist, 1999). The simplifying transformations are necessary for the code generator and are all implemented, except *inlining calls to apply* which is used for partially

applied functions. The optimizing transformations significantly alter the program and achieve similar effects to deforestation (Wadler, 1988) and listlessness (Wadler, 1984). We have only implemented *copy propagation*. This project aims to combine lazy functional programming with precise reference counting. Producing the most optimized code is out of the scope of this project. However, this is something we would like to explore in future research.

### 3 Precise reference counting

After the GRIN transformations, we insert reference counting operations to automatically manage memory. We use an algorithm based on Perceus (Reinking et al., 2021). Perceus is a deterministic syntax-directed algorithm for the linear resource calculus  $\lambda_1$ .  $\lambda_1$  is an untyped lambda calculus extended with explicit binds and pattern matching. Our implementation uses GRIN which have explicit memory operations and different calling conventions. In this section, we give brief a overview of the algorithm and discuss some challenges when adapting Perceus to GRIN. We also describe two optimizations: *drop specialization* and *dup/drop fusion*.

The algorithm uses two sets of resource tokens; an owned environment and a borrowed environment. Elements in the owned environment must be consumed exactly once. We consume a value by calling the function `drop`, or by transfer ownership to another consumer. An example of this is `main` which require no reference counting operations because it transfers ownership of both of its allocations. The allocation `"store (Cnat 100) ; λ x20 →"` is consumed by the suspended computation `"store (FdownFrom x20) ; λ x19 →"`, which in turn is consumed by `"sum x19"`. Elements in the borrowed environment can only be applied to non-consuming operations, such as pattern matching. We can promote an element from the borrowed environment to the owned environment by calling the function `dup`.

Figure 2a presents the function `sum` with reference counting operations inserted and highlighted in gray. Many aspects of adapting Perceus to GRIN are present in `sum`. For example, we can conclude

```

sum x14 =
  fetch x14 [2] ; λ x40 →
  downFrom x40 ; λ x59 x60 x61 →
  case x59 of
    [] →
      update x14 (C[]) ; λ () →
      drop x14 ; λ () →
      unit (Cnat 0)
    _::_ →
      dup x61 ; λ () →
      dup x60 ; λ () →
      update x14 (C_::_ x60 x61) ; λ () →
      drop x14 ; λ () →
      store (Fsum x61) ; λ x10 →
      _+_ x10 x61

```

(a) *dup/drop insertion*

```

sum x14 =
  fetch x14 [2] ; λ x40 →
  downFrom x40 ; λ x59 x60 x61 →
  case x59 of
    [] →
      update x14 (C[]) ; λ () →
      fetch x14 [0] ; λ x499 →
      (case x499 of
        1 → free x14
        - →
          PSub x499 1 ; λ x498 →
          update x14 [0] x498
      ) ; λ () →
      unit (Cnat 0)
    _::_ →
      dup x61 ; λ () →
      dup x60 ; λ () →
      update x14 (C_::_ x60 x61) ; λ () →
      fetch x14 [0] ; λ x501 →
      (case x501 of
        1 →
          drop x61 ; λ () →
          drop x60 ; λ () →
          free x14
        - →
          PSub x501 1 ; λ x500 →
          update x14 [0] x500
      ) ; λ () →
      store (Fsum x61) ; λ x10 →
      _+_ x10 x61

```

(b) *drop specialization*

```

sum x14 =
  fetch x14 [2] ; λ x40 →
  downFrom x40 ; λ x59 x60 x61 →
  case x59 of
    [] →
      update x14 (C[]) ; λ () →
      fetch x14 [0] ; λ x499 →
      (case x499 of
        1 → free x14
        - →
          PSub x499 1 ; λ x498 →
          update x14 [0] x498
      ) ; λ () →
      unit (Cnat 0)
    _::_ →
      update x14 (C_::_ x60 x61) ; λ () →
      fetch x14 [0] ; λ x501 →
      (case x501 of
        1 →
          free x14
        - →
          dup x61 ; λ () →
          dup x60 ; λ () →
          PSub x501 1 ; λ x500 →
          update x14 [0] x500
      ) ; λ () →
      store (Fsum x61) ; λ x10 →
      _+_ x10 x61

```

(c) *dup push-down and dup/drop fusion*

**Figure 2:** *Perceus transformations*

that the operations **fetch** and **update** must be borrowing operations. This is evident because both **fetch** and **update** use `x14` prior to it being explicitly dropped. We can also conclude that **fetch** bind variables that extend the owned environment. Meanwhile, the bound variables of function applications extend the borrowed environment.

One of the goals of GRIN is to improve register utilization for lazy functional languages (Boquist, 1999). As such, the result of function applications, **fetch**, and **unit** are regular values stored in registers. There is no point to reference count values in registers, so we need to treat pointer variables and non-pointer variables differently. This is in contrast to the calculus which Perceus is defined for,  $\lambda_1$ , where all values except variables are heap-allocated (Reinking et al., 2021). Another difference is that functions in GRIN are not allowed to return unboxed pointers. Moreover, all function calls return explicit nodes rather than node variables. For example: `"downFrom x40 ;  $\lambda$  x59 x60 x61  $\rightarrow$  ..."`. This is problematic because `x60` and `x61` are undefined when `downFrom` returns the empty list. Therefore, we are only allowed to dup the variables once the tag is known.

As a result of GRIN’s explicit memory operations, we can describe the Perceus primitives `dup`, `drop`, `is-unique`, `decref`, and `free` with GRIN’s existing constructs. GRIN lacks a primitive for deallocating memory, so `free` is just a foreign function call to `libc`’s `free`. However, our implementation of `drop` is unsatisfactory. Nodes of different tags vary in arity and the arguments can be either boxed or unboxed. Therefore, we need to pattern match on all possible tags to drop the child references. This is similar to the problem with a general `eval` function (Section 2.2). A crucial difference is that `drop` is recursive, so we cannot always eliminate the general function by specialization and inlining. In Figure 2b, we specialize both calls to `drop`. This eliminates all reference counting operations in the branch for the empty list. Then, we push down the `dup` operations and eliminate `dup/drop` pairs. This optimization is called *dup/-drop fusion* and the result is presented in Figure 2c.

We can eliminate the general `drop` function completely for our example program, but this is not always the case. Reinking et al. presents an additional optimization called *reuse analysis*, which performs destructive updates when possible. This transformation is not yet implemented.

## 4 Result

We have implemented a GRIN interpreter and an LLVM code generator to check that our compiler works and that the programs reclaim all the allocated memory. Our implementation of the code generator is simple. GRIN does not demand a specific memory layout for the node values. The only requirements are that tags should be unique and easy to extract (Boquist, 1999). We use an array of four 64-bit integer values for all heap-allocated nodes. The first two address spaces contain the reference count and the tag, respectively. The node arguments occupy the rest of the array. LLVM IR is a statically typed language, while GRIN is untyped. This discrepancy is not an issue because all functions return nodes, and all variables are pointers or integers. Hence, we store the pointers as integer values and convert them to the pointer type (`ptr`) as required. However, we will probably develop a type system for GRIN once we switch to a more sophisticated memory layout and implement the *general unboxing* optimization (Boquist, 1999).

The GRIN interpreter is convenient because it can evaluate all programs after the `eval inlining` transformation. Meanwhile, the code generator only works for programs after the *simplifying transformations* (see Section 2.2). It is also easier to collect information on the program evaluation. For instance, the final version of our example program allocates the following nodes: 101 `Cnat`, 101 `FdownFrom`, 100 `F_`, and 100 `Fsum`. A total of 402 allocations for our tiny example program is excessive. The main culprit is too much laziness. Another contributing factor is that we always allocate new memory instead of reusing previous allocations

## 5 Conclusion and Future Work

In this paper, we combine lazy functional programming with precise reference counting. Our implementation compiles Agda to GRIN and extends GRIN with precise reference counting instructions. Then, we compile GRIN to LLVM and show that the program reclaims all the memory. Currently, our implementation allocates a lot of nodes. In future research, we would like to minimize allocations by implementing the rest of the GRIN optimizing transformations and the Perceus reuse analysis. We are also interested in developing a type system for GRIN and compiling a larger set of Agda programs.

## References

- Boquist, U. (1995). Interprocedural register allocation for lazy functional languages. *Proceedings of the Seventh International Conference on Functional Programming Languages and Computer Architecture*, 270–281. <https://doi.org/10.1145/224164.224215>
- Boquist, U. (1999). Code optimisation techniques for lazy functional languages.
- Boquist, U., & Johnsson, T. (1996). The grin project: A highly optimising back end for lazy functional languages. *Selected Papers from the 8th International Workshop on Implementation of Functional Languages*, 58–84.
- Collins, G. E. (1960). A method for overlapping and erasure of lists. *Commun. ACM*, 3(12), 655–657. <https://doi.org/10.1145/367487.367501>
- Johnsson, T. (1984). Efficient compilation of lazy evaluation. *SIGPLAN Not.*, 19(6), 58–69. <https://doi.org/10.1145/502949.502880>
- Johnsson, T. (1991). Analysing heap contents in a graph reduction intermediate language. In S. L. P. Jones, G. Hutton, & C. K. Holst (Eds.), *Functional programming, glasgow 1990* (pp. 146–171). Springer London.
- Jones, R., & Lins, R. (1996). *Garbage collection: Algorithms for automatic dynamic memory management*. John Wiley & Sons, Inc.
- Pinto, E. (2023). Perceus for ocaml. <https://www.eltonpinto.me/assets/work/mthesis-perceus-for-ocaml.pdf>
- Podlovics, P., Hruska, C., & Péntzes, A. (2021). A modern look at grin, an optimizing functional language back end. *Acta Cybernetica*, 25(4), 847–876. <https://doi.org/10.14232/actacyb.282969>
- Reinking, A., Xie, N., de Moura, L., & Leijen, D. (2021). Perceus: Garbage free reference counting with reuse. *Proceedings of the 42nd ACM SIGPLAN International Conference on Programming Language Design and Implementation*, 96–111. <https://doi.org/10.1145/3453483.3454032>
- Teeuwissen, J. (2023). Reference counting with reuse in roc. <https://studenttheses.uu.nl/handle/20.500.12932/44634>
- Ullrich, S., & de Moura, L. (2021). Counting immutable beans: Reference counting optimized for purely functional programming. *Proceedings of the 31st Symposium on Implementation and Application of Functional Languages*. <https://doi.org/10.1145/3412932.3412935>
- Wadler, P. (1984). Listlessness is better than laziness: Lazy evaluation and garbage collection at compile-time. *Proceedings of the 1984 ACM Symposium on LISP and Functional Programming*, 45–52. <https://doi.org/10.1145/800055.802020>
- Wadler, P. (1988). Deforestation: Transforming programs to eliminate trees. *Proceedings of the Second European Symposium on Programming*, 231–248.